

JRUBY AND RAILS ON GOOGLE APP ENGINE

Amy Unruh

Overview of Talk

- * a **demo**
- * **JRuby** and why you might want to use it
- * Running a **Rails app on JRuby**
 - * using Java libs from your Ruby code
- * Running a Rails app on **Tomcat** with JRuby
- * Running **Rails 2 on Google App Engine/Java**
 - * accessing the App Engine Datastore with the DataMapper adapter
 - * using App Engine services

DEMO

<http://aju-rails.appspot.com/>



JRuby

- * Java implementation of Ruby (jruby.org)
 - * Ruby 1.8.7 compatible (with a few minor differences)
 - * 1.9 compatible soon
- * Core group of developers: Charles Oliver Nutter, Thomas Enebo, Nick Sieger, Ola Bini, Marcin Mielzynski, Bill Dortch, Wayne Meissner, MenTaLguY

Advantages of JRuby

- * Can access Java APIs— both core libs and third-party packages— in your JRuby scripts
- * Can run Rails apps on Java App Servers like Tomcat
- * Can run Rails apps on platforms (like App Engine) that don't support Ruby
- * Outperforms MRI in many cases
 - * 2x to 10x
- * Better threading support, access to JVM tools

JRuby Gems and Libraries

- * Complete implementation of Ruby core/builtin classes, nearly complete implementation of standard libs
- * Any libraries/gems that use native C must be reimplemented in Java
 - * `jruby-openssl`, ...
 - * Some gems build their native library in a platform-independent manner
 - * **mongrel** is one
- * Most 'critical' gems supported in JRuby now

Accessing JRuby

- * Run `jruby -S` to access the JRuby versions of system-level executable commands

- * `jruby -S gem install <gemname>`
 - * `jruby -S rails <appname>`

- * Use 'jruby' to run Ruby programs

- * `jruby script/server`
 - * `jruby a_ruby_script.rb`

- * invoke `jirb` (instead of `irb`)

JRuby on Rails

- * Uses **JDBC** to talk to MySql

- * **jruby -S gem install rails mongrel jdbc-mysql **
activerecord-jdbcmysql-adapter

- * Modify **database.yml**:
adapter: jdbcmysql

- * Then, just pay attention to which implementation of Ruby you are using:

- jruby -S rails <appname>**

- jruby -S rake db:migrate**

- jruby script/server**

Accessing Java APIs from Ruby

```
require 'java'
```

```
import java.net.URL;
```

```
import com.sun.syndication.fetcher.impl.HttpURLFeedFetcher;
```

* Either add jars to the CLASSPATH, or require them.

* `require 'path/to/mycode.jar'`

Accessing Java APIs

```
require 'java'
```

```
import java.net.URL
```

```
import com.sun.syndication.fetcher.impl.HttpURLFeedFetcher
```

```
class PostsController < ApplicationController
```

```
  def index
```

```
    // use java libs to fetch and parse ATOM feed
```

```
    fetcher = HttpURLFeedFetcher.new
```

```
    url = URL.new("http://twitter.com/statuses/user_timeline/newscientist.atom")
```

```
    sf = fetcher.retrieveFeed(url)
```

```
    @feedTitle = sf.getTitle()
```

```
    @feedDescription = sf.getDescription()
```

```
    @entries = sf.getEntries()
```

```
    @posts = Post.all
```

```
    ...
```

```
  end
```

(DEMO: JRUBY ON RAILS)

Running Rails on a Java App Server

- * Thanks to Nick Sieger's Warbler gem, this is straightforward
 - * creates a .war file from a Rails, Merb, or Rack-based application
 - * bundles `jruby-core`, `jruby-stdlib` and `jruby-rack` .jar files
 - * pulls in jar files from `<app-root>/lib`
 - * installs gems into `WEB-INF/gems`
- * `jruby -S gem install warbler`
- * In `<app-root>`: `% warble`
 - * edit `config/warble.rb` to specify the gems to include
 - * to use `jdbc-mysql`, uncomment:

```
config.gems += ["activerecord-jdbcmysql-adapter", "jruby-openssl"]
```
- * Then copy the .war file to (e.g.) `<tomcat>/webapps` to deploy

DEMO: RAILS ON TOMCAT

GOOGLE APP ENGINE/JAVA

Features & Benefits

- * Google App Engine lets you run your web applications on Google's infrastructure.
- * Highly **scalable**, including many scalable **services**
 - * Distributed **Datastore** based on BigTable
- * Don't need to administer stack or data store
- * **Admin console** for app management and statistics
- * Pay as you go (with small quotas free)

Some Limitations & Considerations

- * No access to file system, no threads, no sockets
 - * So, not all third-party Java code will run on App Engine
- * 30-second request timeout
 - * and, no long-running http connections (no streaming)
- * Can be long spin-up time—problematic for apps that aren't heavily used
- * The Datastore is not relational; this can make existing app ports challenging
 - * (but: App Engine for Business)
- * Request/response size and Datastore *entity* size limits (but: Blobstore)
- * Roadmap includes: background servers that can run longer requests, 'warm' or reserved server instances, raised limits on request/response size

App Engine services

- * The Datastore
- * Task Queues and cron jobs
- * Memcache
- * XMPP
- * Google Accounts User API
- * Mail, Images, Blobstore
- * Upcoming/beta: Channel API (push to client), OAuth support, MapReduce service, OpenId

The Datastore

- ✦ * Based on BigTable, distributed and highly scalable, *not relational*
- * Read time is linear in size of **result**, not stored data
- * All *versions* of an app share the same Datastore
- * Constrains the query types allowed in order to maintain linear access time and allow partitioning. E.g.:
 - * no joins (yet)
 - * can't perform multiple inequality tests in a single query
- * Has bulk loading capabilities

Datastore entities

- * The Datastore holds **entities**

- * Entities have a *kind*, and contain *properties*, or named values, from a set of *core value types* (<http://goo.gl/DOPf>)

- * Properties can be multivalued (*property list*)

- * Entities are **schema-less**

- * Two entities of the same kind do not need to have the same properties, or use the same value types for the same properties.

- * so: no schema-definition migrations necessary

Kind	Friend	Friend
(Entity ID portion of) Key	12345	44
firstName	Bob	Joe
lastName	Smith	
tags	colleague, sports	coffee-drinker
emailAddress		joe@joe.com

A Lot More to the Datastore...

- * How its indexes are structured and maintained
 - * Default and custom indexes
 - * Constraints on query construction
- * Entity *groups*
 - * Can define key-based parent/child relationship between entities
 - * Such entities form a group based on ancestor root
 - * *Transactions* may only occur across entity groups

DataMapper and the Datastore

- * ActiveRecord does not work on App Engine
- * **DataMapper** is a Ruby ORM mapper
 - * <http://datamapper.org/>
- * A DataMapper/Datastore adapter has been created
 - * Josh Moore, David Masover, ...
- * Not all DataMapper features are supported by the Datastore integration
 - * Queries unsupported in the Datastore will throw a runtime error
- * If you're adapting an existing app, **Datastore query constraints may require changing your data model**

Defining Models using DataMapper

- * Specify model properties **in the model files**
- * Use Datastore core value types
 - * Of note: `List` (Datastore property lists), `Text` (strings > 500 chars)
- * Use `Serial` for id

Mapper vs ActiveRecord

```
class Account < ActiveRecord::Base
  named_scope :rich,
    :conditions => 'balance >= 1000'
  named_scope :poor,
    :conditions => 'balance < 100'
end
```

```
class Account
  include Mapper::Resource
  property :id, Serial
  property :balance, Integer

  def self.rich
    all(:balance.gte => 1000)
  end
  def self.poor
    all(:balance.lte => 100)
  end
end
```

(example from
David Masover)

Queries in DataMapper

- * Use hashes and builtin comparison methods

- * `books = Book.all(:price.gt => 20, :title => 'My Story')`

- * sort order: `@books = Book.all(:order => [:title.asc])`

- * Illegal for Datastore (multiple inequality tests):

- `books = Book.all(:price.gt => 20, :title.lt => 'A')`

DataMapper Associations

ActiveRecord Terminology

has_many

has_one

belongs_to

DataMapper Terminology

has n

has 1

belongs_to

- * Many-to-many associations currently require special treatment with the Datastore adapter
- * DataMapper docs do not apply

```
class Category
  include DataMapper::Resource

  property :id,      Serial
  property :name, String, :required => true, :length => 500
  timestamps :at
  has n, :books
end
```

Running a Rails App on GAE/J

- * Unofficial, but works pretty well already!
- * Under active development
- * <http://code.google.com/p/appengine-jruby/wiki/RunningRails> ,
<http://groups.google.com/group/appengine-jruby> ...
- * (John Woodell, David Masover, Ryan Brown, ...)

Configuration/Setup for GAE/Rails

* <http://gist.github.com/486250>

* Use MRI, not JRuby (JRuby jars will be bundled)

```
* sudo gem install google-appengine  
  sudo gem install rails -v "2.3.8"  
  sudo gem install rails_dm_datastore  
  sudo gem install activerecord-nullldb-adapter
```

* ...then run a setup script

Configuration for GAE/Rails (cont'd)

- * Files of note generated by the setup script:
 - * `config.ru`
 - * `app.yaml`
 - * is used to auto-generate the `web.xml` and `appengine-web.xml` files
 - * specifies App Engine app id and version
 - * map request handlers, specify authentication requirements
 - * `Gemfile`
 - * uses Gem Bundler; result is `WEB-INF/lib/gems.jar`
 - * `./script .sh` files
 - * `WEB-INF` directory
 - * config file changes (`database.yml`, `environment.rb`, ...)

Example: app.yaml File

application: aju-rails

version: plist3

runtime: jruby

inbound_services:

- xmpp_message
- mail

handlers:

- **url:** /remote_api/*
 login: admin
 servlet: com.google.apphosting.utils.remoteapi.RemoteApiServlet
 name: remoteapi
- **url:** /mailnotif
 login: admin
 servlet: com.example.MailNotifServlet
 name: mailnotif
- **url:** /staff/*
 login: required

Example: generating a scaffold and model

- * Generate a `dd_model`

- * can access Datastore core value types

- * don't need a migration for the 'schema'

- * `./script/generate scaffold post title:string content:text \`
`posted:date tags:list -f --skip-migration`

- * `./script/generate dd_model post title:string content:text \`
`posted:date tags:list -f`

The result: the **Post** model

```
class Post
  include DataMapper::Resource

  property :id,      Serial
  property :title,   String,      :required => true, :length => 500
  property :content, Text,        :required => true, :lazy => false
  property :posted,  Date,        :required => true
  property :tags,    List,        :required => true
  timestamps :at
end
```

Demo: running on App Engine's local Dev Server

- * `./script/server.sh` : generates java app configuration files from `app.yaml`, then starts up App Engine development server
- * `http://localhost:8080/_ah/admin/` : local admin console, access to local 'Datastore'
- * `WEB-INF/appengine-generated/local_db.bin` : holds local Datastore

Deployment to App Engine

- * Create App Engine app id after signing in at appengine.com
- * Define app id and app *version* in app.yaml file
- * `./script/publish.sh`

Using App Engine Services via their Ruby APIs

* Require the service(s), e.g.:

```
* require 'appengine-apis/memcache'  
  require 'appengine-apis/labs/taskqueue'  
  require 'appengine-apis/users'
```

Using the App Engine Memcache

```
def index
  cache = AppEngine::Memcache.new
  @books = cache.get(:books)
  if @books.nil?
    @books = Book.all(:order => [ :title.asc ])
  end
  cache.set(:books, @books)
end

def update
  cache = AppEngine::Memcache.new
  cache.delete(:books)
  ....
end
```


Adding a Task to the App Engine Task Queue

```
prms = {"booktitle" => @book.title}
AppEngine::Labs::TaskQueue.add(:url => "/mailnotif", :params => prms)
```

- * Tasks called via *webhooks*, run asynchronously with admin privs
 - * multiple tasks may be run in parallel
 - * tasks are retried if they fail (return an exception), so should be *idempotent*
- * Multiple task queues may be defined, with different throughputs and 'bucket sizes'
 - * defined in WEB-INF/queue.xml
 - * This example uses the default task queue

Authentication with the App Engine Users API

* app.yaml:

```
- url: /staff/*  
  login: required
```

```
<% if AppEngine::Users.logged_in? %>  
  <%= link_to "Go to staff site", :controller => "staff", :action => "index" %>  
<% else %>  
  <%= link_to "Log in to staff site",  
    AppEngine::Users.create_login_url(url_for(:controller => "staff", :action => "index")) %><br>  
<% end -%>  
  
<%= link_to "Logout", AppEngine::Users.create_logout_url(url_for(:controller => "bookstore",  
  :action => "index")) %>
```

Integrating Servlet Request Handlers with Rails

- * `require 'appengine-rack/java'` in `config.ru`
 - * (However, this did not work for me; I needed to copy the `java.rb` file to my app)
- * Allows direct access to App Engine Services via the Java API
- * Define servlet/URL mapping in `app.yaml`
- * Build jars with App Engine & servlet jars in `CLASSPATH`; move result to `WEB-INF/lib`

```
#!/bin/sh
export CLASSPATH=./build/lib/servlet.jar:./WEB-INF/lib/appengine-api-1.0-sdk-1.3.5.jar

javac com/example/Plain.java
javac com/example/MailNotifServlet.java
mkdir -p WEB-INF/lib
jar -cvf WEB-INF/lib/examples.jar com/example/*.class
rm com/example/*.class
```


Example: A Task to send email & XMPP Notifications

```
public class MailNotifServlet extends HttpServlet {

    public void doPost(HttpServletRequest req, HttpServletResponse resp)
        throws IOException {
        Properties props = new Properties();
        Session session = Session.getDefaultInstance(props, null);
        String booktitle = req.getParameter("booktitle");
        String msgBody = "a book, '" + booktitle + "', has been added...";
        try {
            // first, send XMPP msg
            JID jid = new JID("amygdala@jabber.org");
            logger.info("sending XMPP msg to: " + jid);
            com.google.appengine.api.xmpp.Message xMessage = new MessageBuilder()
                .withRecipientJids(jid)
                .withBody(msgBody)
                .build();
            boolean result = sendMessage(xMessage, jid);
            logger.info("result of sending the XMPP message: " + result);

            // next, email
            javax.mail.Message msg = new MimeMessage(session);
            msg.setFrom(new InternetAddress("aunruh@gmail.com", "The Admin"));
            msg.addRecipient(Message.RecipientType.TO,
                new InternetAddress("amy@infosleuth.net", "Ms. User"));
            msg.setSubject("A book has been added to the catalog.");
            msg.setText(msgBody);
            Transport.send(msg);
        } ...
    }
}
```

Accessing Java Libraries from Rails on App Engine

- * Add third-party jar files to `WEB-INF/lib`
- * Add jars to `CLASSPATH` when running locally
- * Access in your Ruby code as shown
- * Not all third-party Java libraries will run on App Engine

Summary

- * GAE/J + Rails is a promising and robust integration of technologies
 - * Access to App Engine services & automatic scalability is powerful
 - * Rails 3 support in the works
 - * ...just need to get that spin-up-time issue resolved...
- * To keep an eye on: Duby (Mirah)
 - “Mirah is essentially Ruby syntax for writing Java code”
 - <http://github.com/headius/mirah>

END

Many-to-Many Associations using property lists

```
class Author
  include DataMapper::Resource

  property :id,      Serial
  property :name, String, :required => true, :length => 500
  timestamps :at

  has n, :books, :child_key => [:author_ids]
end
```

Many-to-Many Associations using property lists (cont'd)

```
class Book
  include DataMapper::Resource

  property :id,      Serial
  property :title,   String, :required => true, :length => 500
  property :price,   Float,   :required => true
  timestamps :at

  property :author_ids, List
  has n, :order_items
  belongs_to :category

  def authors
    auths = []
    for a in author_ids do
      auths << Author.get(a)
    end
    return auths
  end
end
```


Many-to-Many Associations using an intermediate 'table'

```
class Book
  include DataMapper::Resource

  property :id,      Serial
  property :title,   String,  :required => true, :length => 500
  property :price,   Float,    :required => true
  timestamps :at

  has n, :authorships
  has n, :order_items
  belongs_to :category

  def authors
    auths = []
    for athsh in self.authorships do
      auths << athsh.author
    end
    return auths
  end
end
```

Many-to-Many Associations using an intermediate 'table' (cont'd)

```
class Author
  include DataMapper::Resource

  property :id, Serial
  property :name, String, :required => true, :length => 500
  timestamps :at
  has n, :authorships

  def books
    books = []
    for athsh in self.authorships do
      books << athsh.book
    end
    return books
  end
end
```