

ActiveRecord::Relation and Arel

Rob Mack, Spiceworks

Sunday, September 2, 12

Theo Mills and I built data access layer for a large Rails app. We built it on top of ActiveRecord and Arel. ActiveRecord a bit of a leaky abstraction. You think in associations some places, but table names other places.

What's Arel?

- Query generation for Rails 3
- Relational Algebra (algebra of sets)
- `gem github://rails/arel`

Sunday, September 2, 12

- closure – not that closure, just means returns a set that you can operate on
- creates syntax tree
 - leaves are relations
 - internal nodes are operators

Arel History

- 2007 - Nick Kallen (Pivotal/Twitter) creates `sql_algebra`
- 2008 - Renamed to Active Relation, then Arel
- 2009 - Bryan Helmkamp (@brynary), takes over as maintainer
- 2009 - Emilo Tagua - integrates it with ActiveRecord
- 2010 - Aaron Patterson (@tenderlove) - starts

Source: <https://speakerdeck.com/u/tomstuart/p/relational-algebra-and-arel>

Sunday, September 2, 12

What does this have to do with Rails? Rails 2 queries were generated by string concatenation. Sql string was built up, then executed. Rails 3 changes this with Arel, but when 3.0 launched, Arel was one of the causes for the Rails 3 slowness. The original Arel stores method calls in a link list, to generate the sql, it has to walk backwards through the list. Tenderlove changes design to abstract syntax tree with a manager object. Once you have your AST, you don't necessarily need to generate SQL, hence visitors.

Arel vs. ActiveRecord::Relation

Sunday, September 2, 12

There's a lot of confusion about Arel vs. ActiveRecord::Relation. AR::Rel is the AR layer on top of Arel.

ActiveRecord::Relation

- where
- select
- group
- order
- reorder
- reverse_order
- limit
- offset
- joins
- includes
- lock
- readonly
- having
- uniq

What's in it for me?

- lazy evaluation
- chaining
- portability

Lazy Evaluation

```
# app/controller/user_controller.rb
```

```
def index  
  @recent_users = User.where("created_at > ?", 1.day.ago)  
end
```

```
# app/controller/index.html.erb
```

```
<% cache("recent_users") do %>  
  <% @recent.users.each do |u| %>  
    ...  
  </% %>  
</% %>
```

Chaining

```
> User.where(:id => 1).object_id  
=> 43078
```

```
> User.where(:id => 1).object_id  
=> 43080
```

```
User.where(:status => 'active').  
  where("created_at > ?", 1.day.ago).  
  order(:created_at).limit(10)
```


Useful Methods

```
> User.order(:created_at).to_sql  
=> "SELECT users.* FROM users ORDER BY  
created_at"
```

```
> User.unscoped.class  
=> ActiveRecord::Relation
```

```
> User.unscoped.to_sql  
=> "SELECT users.* FROM users"
```

Merge Relations

```
inactive = User.where(:status => 'inactive')  
recent = User.where("created_at > ?", 1.week.ago)
```

```
recent.merge(inactive).to_sql  
=> "SELECT users.* FROM users WHERE users.status  
= 'active' AND (users.created_at > '2012-08-21');"
```

where_values

```
recent.merge(inactive).where_values  
=> [<Arel::Node>, "created_at > '2012-08-21'"]
```

joins_values
select_values
limit_values

uniq_values
order_values
etc..

A lot like named scopes

```
scope :active, where(:status => 'active')
```

```
def self.active  
  where(:status => 'active')  
end
```

Console Gotcha

```
> User.where(:status => 'active')  
=> doh!
```

```
# File activerecord/lib/active_record/relation.rb  
def inspect  
  to_a.inspect  
end
```

Sunday, September 2, 12

executing this in the rails console is generally bad, use `.to_sql` or `;nil`
Inspect method calls `to_a`, so if you are playing with Relations on console, the output spew can be overwhelming

Arel

Arel Operators

- eq
- not_eq
- lt
- lteq
- gt
- gteq
- matches # like
- not_matches
- in
- not_in
- take # limit
- skip # offset

ActiveRecord::Base

arel_table
arel_engine

u = User.arel_table
u.class
=> Arel::Table

Mix and Match

```
u = User.arel_table
```

```
User.where(u[:created_at].gt(1.day.ago)).limit(1)
```

```
User.where(:id => 10).arel.class
```

```
=> Arel::SelectManager
```

Sunday, September 2, 12

drop arel into relations, use arel method to get from a relation back to the underlying arel object

Why do I need Arel?

- complex queries
- complex dynamic joins, specifying table alias
- compatibility

Visitors

Arel::Visitors::PostgreSQL

Arel::Visitors::MySQL

Arel::Visitors::MSSQL

Arel::Visitors::Oracle

Arel::Visitors::SQLite

Arel::Visitors::IBM_DB

Arel::Visitors::Informix

Sunday, September 2, 12

Visitor design pattern, separate algorithm from object structure that it operates on. Visitors traverse the AST creating SQL.

Portability with Visitors

```
u = User.arel_table  
User.where(u[:name].matches("bob")).to_sql
```

```
"users.name ILIKE 'bob'" # postgres
```

```
"users.name LIKE 'bob'" #mysql
```

Table Join Gotcha

```
def self.cheaper_than(price)
  where("price < ?", price)
end
```

```
def self.cheaper_than(price)
  where(arel_table[:price].lt(price))
end
```

Sunday, September 2, 12

second one will work if you are joining in multiple tables with price column, aliased table or otherwise

Ernie Miller

- `github://ernie/squeel`
- railscasts.com/episodes/354-squeel
- erniemiller.org

Sunday, September 2, 12

Ernie Miller – living social, knows a ton about arel, contributor, interesting blog
before you go diving to far into Arel, take a look at the squeel gem


Links

- Aaron Patterson on Arel 2.0 - <http://web.archive.org/web/20101113093529/http://engineering.attinteractive.com/2010/10/arel-two-point-ohhhhh-yaaaaaa/>
- History of Arel - <https://speakerdeck.com/u/tomstuart/p/relational-algebra-and-arel>

Thanks!

Rob Mack, Spiceworks
@robmack
rob@robmack.com

Obligatory “We’re Hiring!” Slide



SPICEWORKS™
IT'S EVERYTHING IT